



INTERNSHIP NOTES

| | |
|-----------|----------------------------|
| ⌚ Created | @October 22, 2025 11:04 PM |
| ⌚ Class | IIITH |
| ☰ Event | |

The core concepts and functions used in the Python code blocks, focusing on the `ultralytics` library.

Task1: Basic image segmentation and classification

TASK 1: Code

Explanation of the YOLOv8 Python Code

The core of the procedure relies on the `ultralytics` Python package, which provides a simple, high-level API for using the powerful YOLO models.

Core Components

The two essential lines that appear in most steps are:

1. `from ultralytics import YOLO`
2. `model = YOLO("model_file.pt")`

- `from ultralytics import YOLO` : This imports the main `YOLO` class from the installed library. This class is the central object you'll interact with for loading models, running predictions, and accessing results.
- `model = YOLO("yolov8n.pt")` : This **loads the pre-trained model**.
 - `YOLO()` : Creates an instance of the YOLO object.

- "yolov8n.pt" : This is the model file. YOLOv8 comes in different sizes (nano 'n', small 's', medium 'm', etc.). The 'n' stands for **nano**, the smallest and fastest version, which is great for quick testing. The file is downloaded automatically the first time you run the script.
 - For **detection**, you use `yolov8n.pt` .
 - For **segmentation** (Step 6), you use `yolov8n-seg.pt` , which is a segmentation-specific version.
-

Explanation of the `model.predict()` Function

The `model.predict()` function is the single, unified way to run the model on data (images, videos, or streams). Here is a breakdown of the arguments used:

results = model.predict(...)

| Argument | Example Value | Purpose |
|----------------------|---|---|
| <code>source</code> | "https://..." or <code>r"C:\..."</code> | Mandatory: Specifies the input to process. This can be a local file path (image/video), a folder path (to process all images), or a URL . |
| <code>show</code> | <code>True</code> or <code>False</code> | If <code>True</code> , a window will pop up to display the results in real-time as the model processes the image/video. Useful for live preview. |
| <code>save</code> | <code>True</code> or <code>False</code> | If <code>True</code> , the annotated image (with bounding boxes or masks) is saved to the output directory. |
| <code>project</code> | <code>r"C:\Users\...\YOLO_Outputs"</code> | Optional: The root directory where all prediction runs are saved. If not specified, it defaults to a folder like <code>runs/detect</code> or <code>runs/segment</code> in your current working directory. |
| <code>name</code> | "Paypal_Mafia_2014" | Optional: The specific folder name for this particular run, created inside the <code>project</code> directory. This keeps your experiments organized. |
| <code>classes</code> | <code>[0]</code> | Optional/Crucial for Filtering: This is a list of class IDs you want the model to process or |

display. It's used to filter results.

- YOLO is trained on the **COCO dataset**, where **class ID `0` always corresponds to a "person"**.
* By setting `classes=[0]`, you instruct the model to only output results (bounding boxes or segmentation masks) for detected persons, ignoring all other objects (cars, bikes, animals, etc.). |

The `results` Variable

The line `results = model.predict(...)` captures the output.

- The `results` variable is a list of `Results` objects. For a single image prediction, it usually contains one `Results` object.
- This object contains all the detailed, low-level information about the prediction, such as:
 - The coordinates of all bounding boxes.
 - The confidence score for each detection.
 - The segmentation masks (for segmentation models).
 - The detected class IDs.

The Key Difference: Detection vs. Segmentation

The main difference between Step 5 and Step 6 is which model is loaded:

| Feature | Step 5: Object Detection | Step 6: Instance Segmentation |
|------------|--|--|
| Goal | Find objects and draw a bounding box around them. | Find objects and generate a pixel-level mask for the exact shape of the object. |
| Model File | <code>yolov8n.pt</code> | <code>yolov8n-seg.pt</code> |
| Output | Rectangles/Boxes around people. | Pixel-perfect shapes (masks) highlighting the entire area of each person. |

By using the `yolov8n-seg.pt` model and the `classes=[0]` filter, you can run an , person-only instance segmentation task on your local image.

Ultralytics YOLOv8 library to perform two distinct tasks on a folder of images

Task 2 Code (25-10-2025)

This code block demonstrates how to use the **Ultralytics YOLOv8 library** to perform two distinct tasks on a **folder of images**

Precision, Recall, and F1 Score are essential metrics for evaluating how well a model performs **classification** or **detection** tasks. They are all derived from the fundamental counts of correct and incorrect predictions: **True Positives (TP)**, **False Positives (FP)**, and **False Negatives (FN)**.

| Term | Definition | Outcome |
|----------------------------|--|--|
| True Positive (TP) | The model correctly identifies an object. | A person is present, and the model drew a bounding box around them. |
| False Positive (FP) | The model incorrectly identifies something as an object. | The model drew a bounding box around a tree, claiming it's a person. |
| False Negative (FN) | The model misses an object that is present. | A person is present, but the model failed to draw a bounding box. |

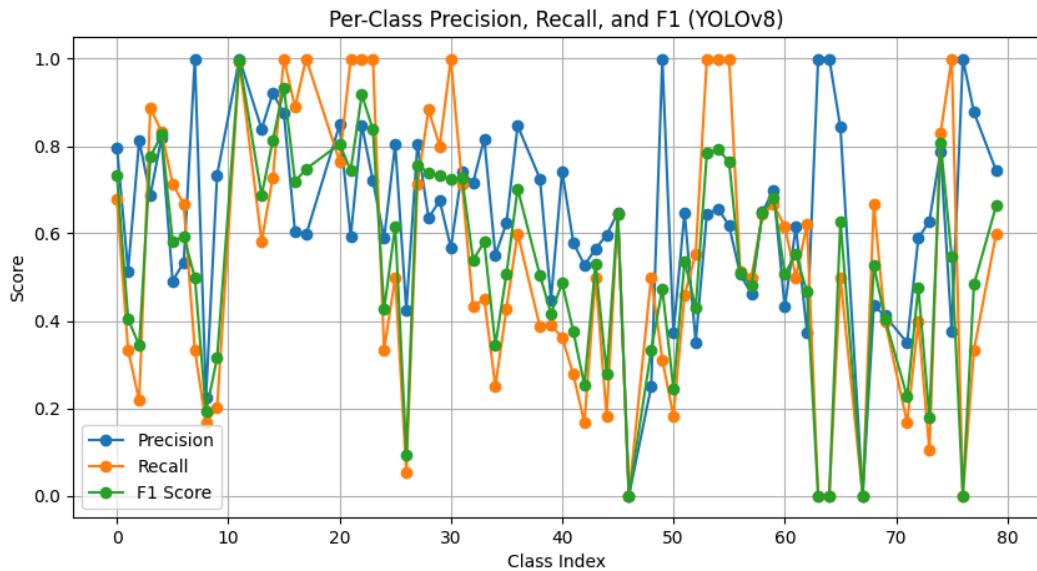
(Note: True Negative (TN) is typically ignored in object detection because the vast majority of an image is "negative" (empty space).) **because these metrics are designed to focus on the performance of the positive class**

`metrics = model_detect.val(data='coco128.yaml')` : The `.val()` function runs the loaded model against a **validation dataset** (here, the small COCO128 dataset is used as an example). This function calculates performance metrics.

`metrics.box.p` / `metrics.box.r` / `metrics.box.f1` : These access the **Per-Class Metrics**.
• **Precision (\$\mathbf{P}\$)**: Of all objects the model *detected* for a class, how

many were correct? (High P = low false positives).

- **Recall (\mathbf{R}):** Of all actual objects in a class, how many did the model detect? (High R = low false negatives).
- **$\mathbf{F1}$ Score:** The harmonic mean of Precision and Recall.



YOLOv8 Performance Metrics (Per-Class Precision, Recall & F1 Score)

This graph shows how accurately the model detects each object class:

- **Precision (blue):** How often the detected objects are correct.
- **Recall (orange):** How well the model finds all instances of an object.
- **F1 Score (green):** The balance between Precision and Recall.

Each point represents a class (like person, car, chair, etc.), showing how YOLOv8 performs across all detected objects

TASK 3:

```
from ultralytics import YOLO  
import os
```

```
# Load YOLO model
model = YOLO('yolov8n.pt') # You can use yolov8m.pt or yolov9.pt too

# Input/output directories
input_folder = r"C:\Users\aryan\OneDrive\Desktop\frames"
output_folder = r"C:\Users\aryan\OneDrive\Desktop\frames_detected"

os.makedirs(output_folder, exist_ok=True)

# Process all frames
results = model.predict(source=input_folder, save=True, project=output_folder, name='')
```

```
from ultralytics import YOLO
import os

# Load YOLO segmentation model
model_seg = YOLO('yolov8n-seg.pt')

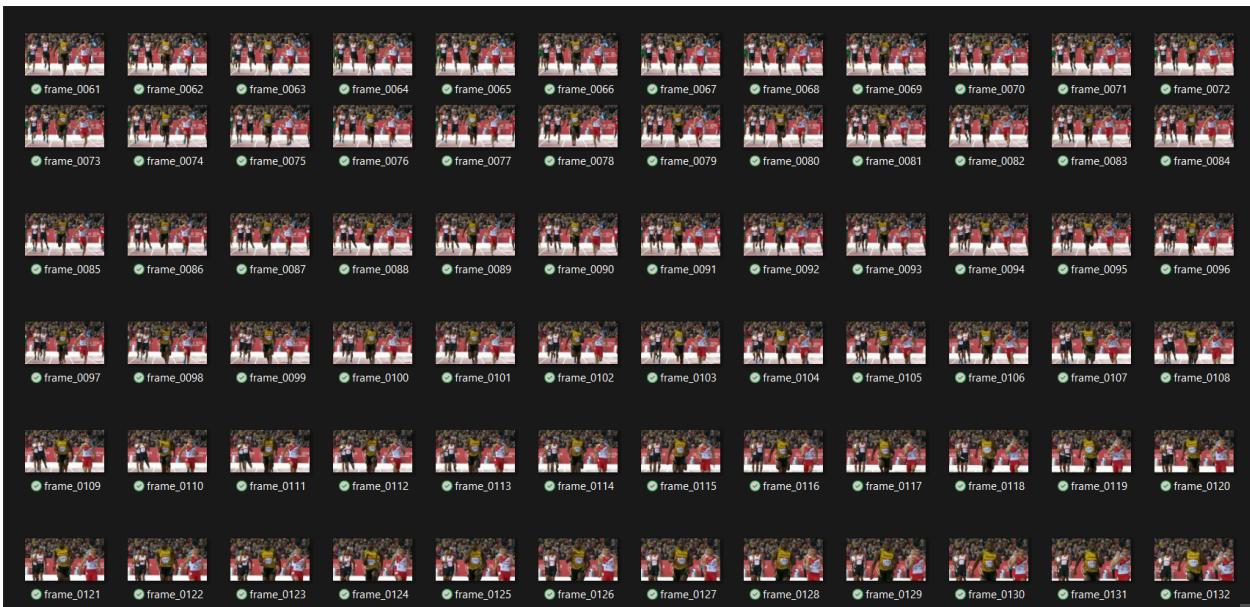
input_folder = r"C:\Users\aryan\OneDrive\Desktop\frames"
output_folder_seg = r"C:\Users\aryan\OneDrive\Desktop\frames_segmented"

os.makedirs(output_folder_seg, exist_ok=True)

results = model_seg.predict(source=input_folder, save=True, project=output_folder_seg, name='')
```



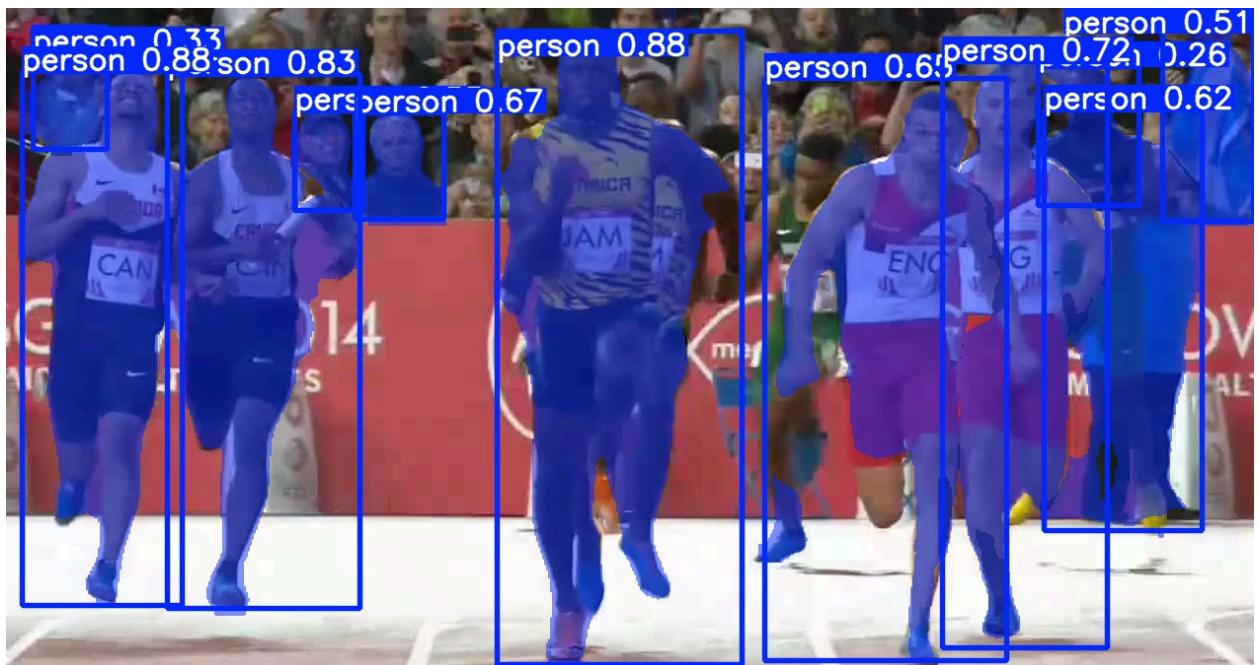
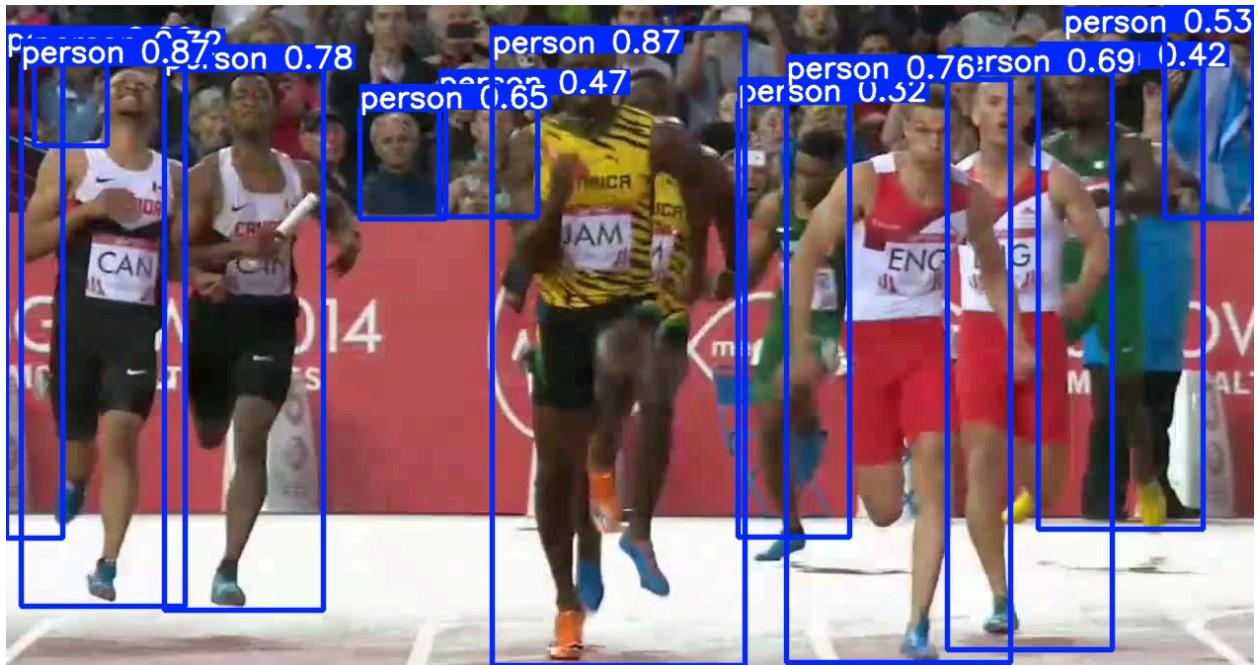
the original video:
using FFmpeg , i have divided the video into 132 Frames.



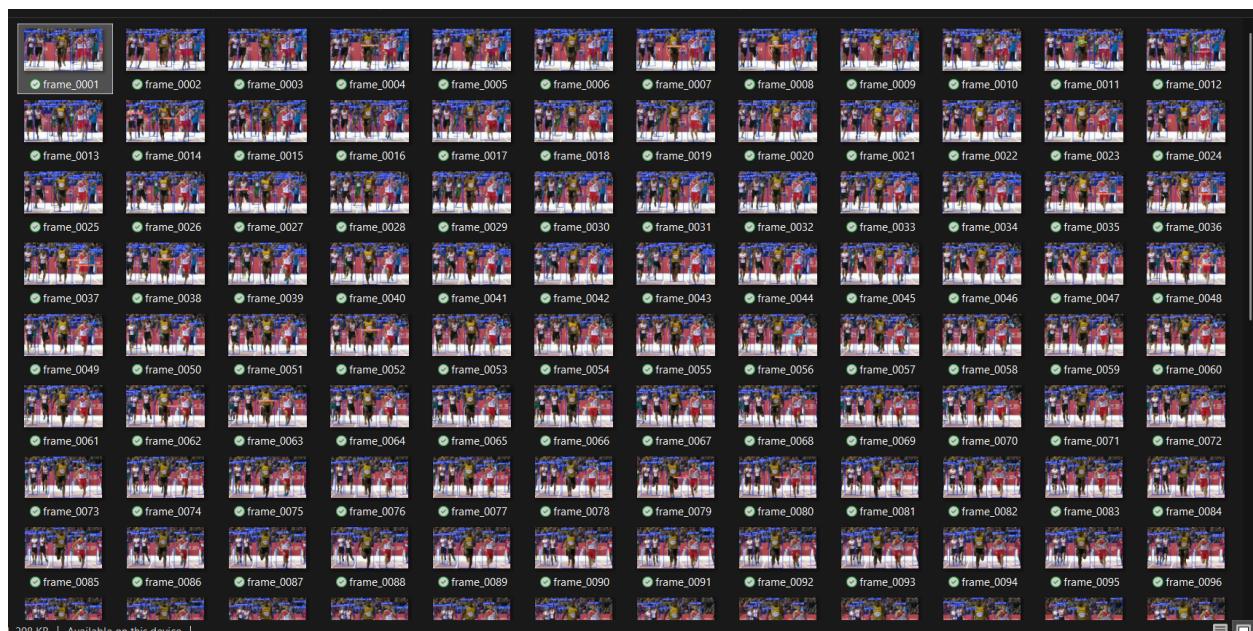
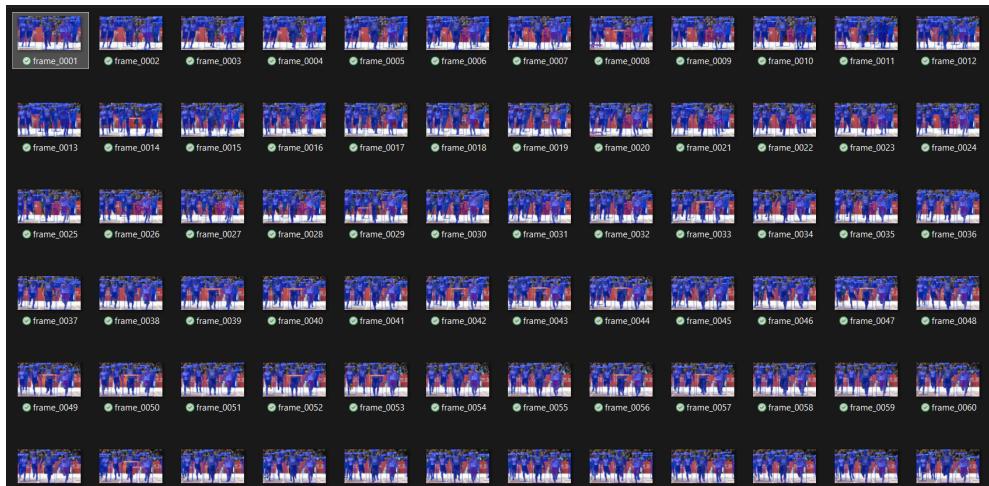
Saved it under Frames folder

then, i have applied Multiple image Object Detection and Segmentation, as learned in TASK 2.

These are the Sample Detected and segmented Images ,



Similarly the segmentation and detection is applied to all the 132 frames of the video



Rebuilding Videos from Frames.



Stitching Images into a Video (Using FFmpeg)

After performing object detection or segmentation on each frame, you generated hundreds of processed images like:

```
frame_0001.jpg  
frame_0002.jpg  
frame_0003.jpg  
...
```

These were saved inside:

- For **Detection**:

```
C:\Users\aryan\OneDrive\Desktop\frames_detected\predict
```

- For **Segmentation**:

```
C:\Users\aryan\OneDrive\Desktop\frames_segmented\predict
```

Step 1: Convert Processed Frames → Video

You used **FFmpeg** (installed manually) to combine those frames back into a video.

Command for Detection Video:

```
& "C:\Users\aryan\Downloads\ffmpeg-2025-10-27-git-68152978b5-full_build  
\ffmpeg-2025-10-27-git-68152978b5-full_build\bin\ffmpeg.exe" -framerate 10  
-i "C:\Users\aryan\OneDrive\Desktop\frames_detected\predict\frame_%04d.jp  
g" -c:v libx264 -pix_fmt yuv420p "C:\Users\aryan\OneDrive\Desktop\detected  
_video.mp4"
```

Command for Segmentation Video:

```
& "C:\Users\aryan\Downloads\ffmpeg-2025-10-27-git-68152978b5-full_build  
\ffmpeg-2025-10-27-git-68152978b5-full_build\bin\ffmpeg.exe" -framerate 10  
-i "C:\Users\aryan\OneDrive\Desktop\frames_segmented\predict\frame_%04  
d.jpg" -c:v libx264 -pix_fmt yuv420p "C:\Users\aryan\OneDrive\Desktop\seg  
mented_video.mp4"
```

Explanation of the Command

| Parameter | Meaning |
|---------------------|--|
| -framerate 10 | Sets frame rate (10 images per second). You can use 24 for smoother video. |
| -i "frame_%04d.jpg" | Reads images named as frame_0001.jpg, frame_0002.jpg, etc. |
| -c:v libx264 | Uses H.264 codec — standard for MP4 videos. |
| -pix_fmt yuv420p | Ensures compatibility with all video players. |
| "output.mp4" | Name and path of the generated video file. |

Step 2: Merge Multiple Videos Side-by-Side

Once both videos (detection + segmentation) were created, you merged them horizontally for comparison.

Command for Merging Detection + Segmentation:

```
& "C:\Users\aryan\Downloads\ffmpeg-2025-10-27-git-68152978b5-full_build\ffmpeg-2025-10-27-git-68152978b5-full_build\bin\ffmpeg.exe" -i "C:\Users\aryan\OneDrive\Desktop\detected_video.mp4" -i "C:\Users\aryan\OneDrive\Desktop\segmented_video.mp4" -filter_complex "hstack=inputs=2" "C:\Users\aryan\OneDrive\Desktop\final_combined.mp4"
```

Or if you also want to include the **original video** side by side:

```
& "C:\Users\aryan\Downloads\ffmpeg-2025-10-27-git-68152978b5-full_build\ffmpeg-2025-10-27-git-68152978b5-full_build\bin\ffmpeg.exe" -i "C:\Users\aryan\OneDrive\Desktop\input.mp4" -i "C:\Users\aryan\OneDrive\Desktop\detected_video.mp4" -i "C:\Users\aryan\OneDrive\Desktop\segmented_video.mp4" -filter_complex "hstack=inputs=3" "C:\Users\aryan\OneDrive\Desktop\final_combined.mp4"
```

Explanation of Merge Command

| Part | Description |
|--|---|
| <code>-i "file1.mp4" -i "file2.mp4"</code> | Input videos to combine. |
| <code>-filter_complex "hstack=inputs=2"</code> | Combines videos horizontally (use <code>vstack</code> for vertical). |
| <code>"final_combined.mp4"</code> | Output merged video file. |

📁 Folder structure summary:

```

Desktop/
|
|   input.mp4
|   frames/
|       └── frame_0001.jpg ...
|
|   frames_detected/
|       └── predict/
|           └── frame_0001.jpg ...
|
|   frames_segmented/
|       └── predict/
|           └── frame_0001.jpg ...
|
|   detected_video.mp4
|   segmented_video.mp4
|   final_combined.mp4

```

🎯 Final Output:

-  `detected_video.mp4` → bounding boxes drawn
-  `segmented_video.mp4` → colored object masks
-  `final_combined.mp4` → side-by-side comparison video