



INTERNSHIP NOTES

🕒 Created	@October 22, 2025 11:04 PM
📁 Class	IIITH
☰ Event	

The core concepts and functions used in the Python code blocks, focusing on the `ultralytics` library.

Task1: Basic image segmentation and classification

TASK 1: Code

Explanation of the YOLOv8 Python Code

The core of the procedure relies on the `ultralytics` Python package, which provides a simple, high-level API for using the powerful YOLO models.

Core Components

The two essential lines that appear in most steps are:

1. `from ultralytics import YOLO`
 2. `model = YOLO("model_file.pt")`
- `from ultralytics import YOLO` : This imports the main `YOLO` **class** from the installed library. This class is the central object you'll interact with for loading models, running predictions, and accessing results.
 - `model = YOLO("yolov8n.pt")` : This **loads the pre-trained model**.
 - `YOLO()` : Creates an instance of the YOLO object.

- `"yolov8n.pt"` : This is the model file. YOLOv8 comes in different sizes (nano 'n', small 's', medium 'm', etc.). The 'n' stands for **nano**, the smallest and fastest version, which is great for quick testing. The file is downloaded automatically the first time you run the script.
 - For **detection**, you use `yolov8n.pt` .
 - For **segmentation** (Step 6), you use `yolov8n-seg.pt` , which is a segmentation-specific version.

Explanation of the `model.predict()` Function

The `model.predict()` function is the single, unified way to run the model on data (images, videos, or streams). Here is a breakdown of the arguments used:

`results = model.predict(...)`

Argument	Example Value	Purpose
<code>source</code>	<code>"https://..."</code> or <code>r"C:\..."</code>	Mandatory: Specifies the input to process. This can be a local file path (image/video), a folder path (to process all images), or a URL .
<code>show</code>	<code>True</code> or <code>False</code>	If <code>True</code> , a window will pop up to display the results in real-time as the model processes the image/video. Useful for live preview.
<code>save</code>	<code>True</code> or <code>False</code>	If <code>True</code> , the annotated image (with bounding boxes or masks) is saved to the output directory.
<code>project</code>	<code>r"C:\Users\...\YOLO_Outputs"</code>	Optional: The root directory where all prediction runs are saved. If not specified, it defaults to a folder like <code>runs/detect</code> or <code>runs/segment</code> in your current working directory.
<code>name</code>	<code>"Paypal_Mafia_2014"</code>	Optional: The specific folder name for this particular run, created inside the <code>project</code> directory. This keeps your experiments organized.
<code>classes</code>	<code>[0]</code>	Optional/Crucial for Filtering: This is a list of class IDs you want the model to process or

	display. It's used to filter results.
--	---------------------------------------

- YOLO is trained on the **COCO dataset**, where **class ID `0`** always corresponds to a **"person"**.
 * By setting `classes=[0]`, you instruct the model to only output results (bounding boxes or segmentation masks) for detected persons, ignoring all other objects (cars, bikes, animals, etc.). |

The **results** Variable

The line `results = model.predict(...)` captures the output.

- The **results** variable is a list of **Results objects**. For a single image prediction, it usually contains one **Results** object.
- This object contains all the detailed, low-level information about the prediction, such as:
 - The coordinates of all bounding boxes.
 - The confidence score for each detection.
 - The segmentation masks (for segmentation models).
 - The detected class IDs.

The Key Difference: Detection vs. Segmentation

The main difference between Step 5 and Step 6 is which model is loaded:

Feature	Step 5: Object Detection	Step 6: Instance Segmentation
Goal	Find objects and draw a bounding box around them.	Find objects and generate a pixel-level mask for the exact shape of the object.
Model File	<code>yolov8n.pt</code>	<code>yolov8n-seg.pt</code>
Output	Rectangles/Boxes around people.	Pixel-perfect shapes (masks) highlighting the entire area of each person.

By using the `yolov8n-seg.pt` model and the `classes=[0]` filter, you can run an , person-only instance segmentation task on your local image.

Ultralytics YOLOv8 library to perform two distinct tasks on a folder of images

Task 2 Code (25-10-2025)

This code block demonstrates how to use the **Ultralytics YOLOv8 library** to perform two distinct tasks on a **folder of images**

Precision, Recall, and F1 Score are essential metrics for evaluating how well a model performs **classification** or **detection** tasks. They are all derived from the fundamental counts of correct and incorrect predictions: **True Positives (TP)**, **False Positives (FP)**, and **False Negatives (FN)**.

Term	Definition	Outcome
True Positive (TP)	The model correctly identifies an object.	A person is present, and the model drew a bounding box around them. ✓
False Positive (FP)	The model incorrectly identifies something as an object.	The model drew a bounding box around a tree, claiming it's a person. ✗
False Negative (FN)	The model misses an object that is present.	A person is present, but the model failed to draw a bounding box. 🚫

(Note: True Negative (TN) is typically ignored in object detection because the vast majority of an image is "negative" (empty space).) **because these metrics are designed to focus on the performance of the positive class**

`metrics = model_detect.val(data='coco128.yaml')` : The `.val()` function runs the loaded model against a **validation dataset** (here, the small COCO128 dataset is used as an example). This function calculates performance metrics.


`metrics.box.p` / `metrics.box.r` / `metrics.box.f1` : These access the **Per-Class Metrics**.

- **Precision (\mathbf{P})**: Of all objects the model *detected* for a class, how

many were correct? (High P = low false positives).

- **Recall (\mathbf{R})**: Of all actual objects in a class, how many did the model *detect*? (High R = low false negatives).

- **F1 Score**: The harmonic mean of Precision and Recall.



YOLOv8 Performance Metrics (Per-Class Precision, Recall & F1 Score)

This graph shows how accurately the model detects each object class:

- **Precision (blue)**: How often the detected objects are correct.
- **Recall (orange)**: How well the model finds all instances of an object.
- **F1 Score (green)**: The balance between Precision and Recall.

Each point represents a class (like person, car, chair, etc.), showing how YOLOv8 performs across all detected objects